

# HW4

May 27, 2017

```
!pip3 install seaborn
```

I couldn't execute all the utility programs you provided in HW4. So I went ahead and loaded the Bike Share data from pronto-data website.

```
In [4]: !pip install seaborn
        !pip install python-utils
```

```
Requirement already satisfied: seaborn in ./anaconda/lib/python3.5/site-packages
Collecting python-utils
  Downloading python_utils-2.1.0-py2.py3-none-any.whl
Requirement already satisfied: six in ./anaconda/lib/python3.5/site-packages (from
Installing collected packages: python-utils
Successfully installed python-utils-2.1.0
```

```
In [1]: !curl -O https://s3.amazonaws.com/pronto-data/open_data_year_two.zip
        !unzip open_data_year_two.zip
```

```
% Total      % Received % Xferd  Average Speed   Time    Time       Time  Current
                               Dload  Upload   Total   Spent    Left   Speed
100 5627k  100 5627k    0     0  1892k      0  0:00:02  0:00:02 --:--:-- 1993k
Archive:  open_data_year_two.zip
replace 2016_trip_data.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: ^C
```

```
In [2]: import seaborn as sns
        import matplotlib.pyplot as plt
        %matplotlib inline
        import pandas as pd
        bikes = pd.read_csv('2016_trip_data.csv', parse_dates=['starttime', 'stoptime'],
                           infer_datetime_format=True)

        bikes['start'] = pd.to_datetime(bikes['starttime'], infer_datetime_format=True)
        bikes['end'] = pd.to_datetime(bikes['stoptime'], infer_datetime_format=True)
        bikes['hour_of_day'] = (bikes.start.dt.hour + (bikes.start.dt.minute/60)).round()
        bikes.head()

        hours = bikes.groupby('hour_of_day').agg('count')
```


```
hours['hour'] = hours.index
```


```
hours.start.plot()
```

```
hours.start.plot()
```

```
sns.lmplot(x='hour', y='start', data=hours, aspect=1.5, scatter_kws={'alpha
```

Out[2]: <seaborn.axisgrid.FacetGrid at 0x1110a52e8>





1. Create 3 models fit to hour\_of\_day with varying polynomial degrees

```
In [3]: from sklearn.preprocessing import PolynomialFeatures
        from sklearn import linear_model
        import numpy as np
        poly = PolynomialFeatures(degree=2)
        x = np.array(hours['hour']).reshape(-1,1)
        y = np.array(hours['start'])
        x2 = poly.fit_transform(x)
        M1 = linear_model.LinearRegression()
        M1.fit(x2, y)
        (M1.intercept_, M1.coef_)

        poly = PolynomialFeatures(degree=4)


        x4 = poly.fit_transform(x)
        M2 = linear_model.LinearRegression()
        M2.fit(x4, y)
        (M2.intercept_, M2.coef_)

        poly = PolynomialFeatures(degree=10)
        x10 = poly.fit_transform(x)
        M3 = linear_model.LinearRegression()
        M3.fit(x10, y)
```

```
(M3.intercept_,M3.coef_)

plt.scatter (x,y)
plt.plot(x,np.dot(x2,M1.coef_)+M1.intercept_,c='b')
plt.plot(x,np.dot(x4,M2.coef_)+M2.intercept_,c='r')
plt.plot(x,np.dot(x10,M3.coef_)+M3.intercept_,c='y')
```

Out[3]: [<matplotlib.lines.Line2D at 0x1129b29b0>]



Choose one of the polynomial models and create 3 new models fit to hour\_of\_day with different Ridge Regression  $\alpha$  (alpha) Ridge Coefficient values

```
In [4]: M1r = linear_model.Ridge(alpha = 0.1)
M1r.fit(x2, y)
(M1r.intercept_,M1r.coef_)

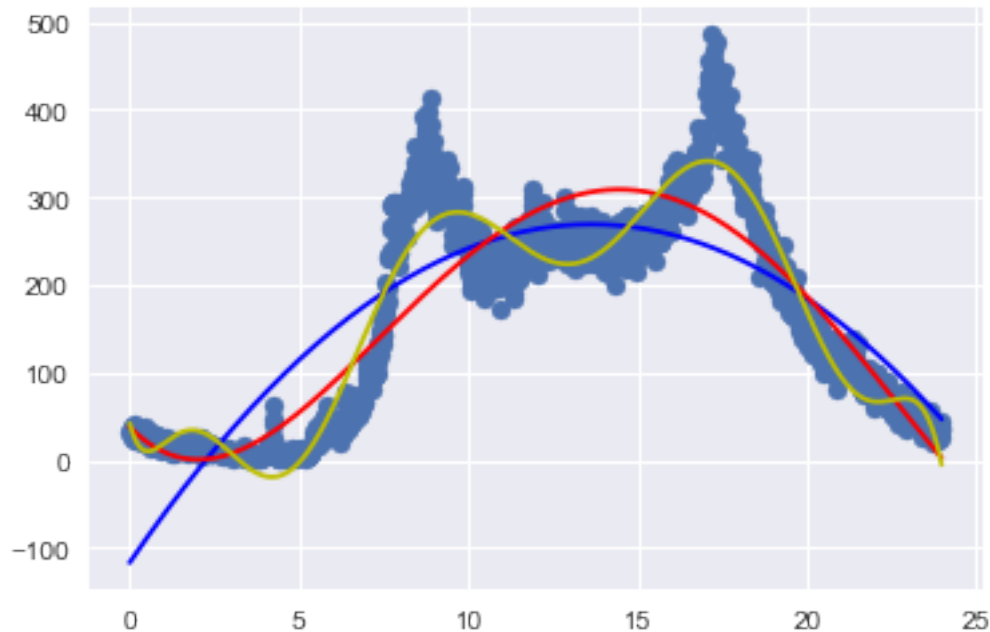
poly = PolynomialFeatures(degree=4)
M2r = linear_model.Ridge(alpha = 0.5)
M2r.fit(x4, y)
(M2r.intercept_,M2r.coef_)

M3r = linear_model.Ridge(alpha = 0.6)
M3r.fit(x10, y)
(M3r.intercept_,M3r.coef_)

plt.scatter (x,y)
```

```
plt.plot(x,np.dot(x2,M1r.coef_)+M1r.intercept_,c='b')
plt.plot(x,np.dot(x4,M2r.coef_)+M2r.intercept_,c='r')
plt.plot(x,np.dot(x10,M3r.coef_)+M3r.intercept_,c='y')
```

Out[4]: [<matplotlib.lines.Line2D at 0x112a35f98>]



```
In [10]: from sklearn import cross_validation
import numpy as np
from sklearn.metrics import mean_squared_error
alphas = 10**np.linspace(10,-2,100)*0.5
X_train, X_test , y_train, y_test = cross_validation.train_test_split(x10,
ridgecv = linear_model.RidgeCV()
ridgecv.fit(X_train, y_train)
print(ridgecv.alpha_)
M3r = linear_model.Ridge(alpha =ridgecv.alpha_ )
M3r.fit(x10, y)
print( (M3r.intercept_,M3r.coef_))

plt.scatter (x,y)
plt.plot(x,np.dot(x10,M3r.coef_)+M3r.intercept_,c='y')


pred2 = M3r.predict(X_test)
print(mean_squared_error(y_test, pred2))
```

```
0.1
(73.653874953652519, array([ 0.00000000e+00, -2.59842321e+02,  3.25978778e+02,
```

```

-1.70658367e+02,  4.50080432e+01,  -6.65883267e+00,
 5.88657968e-01, -3.17775732e-02,  1.02632952e-03,
-1.82067736e-05,  1.36343847e-07]])
1281.73694429

```



```

In [12]: from sklearn import cross_validation
import numpy as np
from sklearn.metrics import mean_squared_error
alphas = 10**np.linspace(10,-2,100)*0.5
X_train, X_test , y_train, y_test = cross_validation.train_test_split(x4,
ridgecv = linear_model.RidgeCV()
ridgecv.fit(X_train, y_train)
print(ridgecv.alpha_)
M2r = linear_model.Ridge(alpha =ridgecv.alpha_ )
M2r.fit(x4, y)
print( (M2r.intercept_,M2r.coef_))

plt.scatter (x,y)
plt.plot(x,np.dot(x4,M2r.coef_)+M2r.intercept_,c='y')

pred1 = M2r.predict(X_test)
print(mean_squared_error(y_test, pred1))

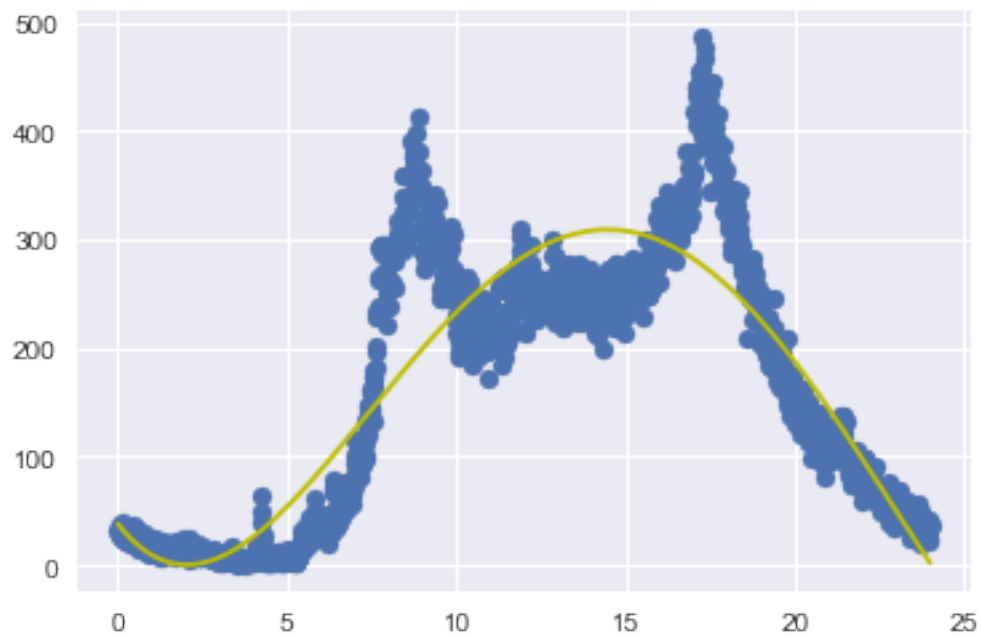
```

```

0.1
(38.144887182994069, array([ 0.00000000e+00, -3.96226930e+01,  1.18274210e+01,

```

```
-7.08852053e-01, 1.17575804e-02]))  
3395.06064824
```



Based on the MSE, I favor a model with 10 polynomial degree, although the model with 4 degrees is parsimonious

In [ ]: